

TRANSFORMING AMERICAN HEALTHCARE

Part IV: Tools for Transformation

Chapters 13–15

PART IV: TOOLS FOR TRANSFORMATION

The preceding chapters described what healthcare transformation requires: the policy environment, the economic framework, the technology infrastructure, the clinical models, the equity design, and the operational architecture. This section describes the analytical tools that make those requirements actionable — the instruments by which planners, executives, clinicians, and policymakers can turn analysis into decision.

Health Transformation Review's platform includes 19 interactive research tools, a structured learning academy, and an AI-powered analyst powered by retrieval-augmented generation (RAG). Each is described in this section as both a practical instrument and an educational framework for the underlying methodologies.

CHAPTER 13: Interactive Analytics — The Research Lab's 19 Tools

The Research Lab Philosophy

The Research Lab is built on a specific premise: the difference between knowing a concept and being able to apply it is practice with realistic data and feedback. A textbook can describe the mathematics of an MSSP shared savings calculation. A calculator that lets an analyst enter their specific population's data and see the resulting shared savings — varying assumptions, seeing sensitivity ranges, understanding break-even scenarios — builds a different kind of competency.

Each of the 19 Research Lab tools is designed as a bridge between conceptual knowledge and applied decision-making. They are organized into six thematic sections:

RESEARCH LAB STRUCTURE
Section 1: INTEROPERABILITY & RISK
Tool 1: FHIR Interoperability Lab
Tool 2: Risk Stratification Engine
Section 2: PAYMENT MODELS & VBC
Tool 3: APM Design Lab
Tool 4: APM Shared Savings Calculator
Tool 5: Cost-Effectiveness Analysis Calculator
Section 3: POPULATION & EQUITY
Tool 6: Population Health Modeler
Tool 7: Health Equity Studio
Section 4: POLICY & QUALITY SCIENCES
Tool 8: Policy Simulator
Tool 9: Clinical Quality Optimizer
Tool 10: Hospital Financial Scorecard
Tool 11: HTA Studio
Tool 12: Actuarial Lab
Section 5: TECHNOLOGY & AI
Tool 13: AI Analytics Lab
Tool 14: Digital Health Lab

Section 6: KNOWLEDGE & WORKSPACE

Tool 15: Evidence Library

Tool 16: Workforce Modeler

Tool 17: Innovation Leaderboard

Tool 18: Research Workspace

Section 1: Interoperability & Risk

Tool 1: FHIR Interoperability Lab

The FHIR Interoperability Lab provides hands-on experience with FHIR R4 resource construction, terminology mapping, CDS Hooks integration, and ONC compliance.

Core capabilities:

FHIR R4 Resource Builder: Construct FHIR resources for 15+ resource types (Patient, Condition, Observation, MedicationRequest, etc.) with real-time schema validation against the HL7 R4 and US Core Implementation Guide profiles. The builder highlights conformance errors and explains the correction required.

Terminology Mapping: Clinical terminologies are the vocabulary of health data exchange. The lab supports:

- SNOMED CT ↔ ICD-10-CM mapping (diagnosis codes)
- LOINC code lookup (lab tests, vital signs)
- CPT ↔ SNOMED procedure mapping
- RxNorm drug code lookup

Understanding terminology mapping is essential because interoperability requires not just data exchange but shared meaning. Two systems that both say "patient has diabetes" but use different codes cannot be reliably aggregated — which is why FHIR mandates specific terminology bindings for critical data elements.

CDS Hooks Simulator: CDS Hooks is the FHIR-based standard for delivering clinical decision support at the point of care within the EHR workflow. The simulator allows configuration of:

- Hook type: **patient-view**, **order-select**, **order-sign**, **encounter-start**
- Context payload: the clinical context (current patient, current order) sent to the CDS service
- Card response: the recommendation returned to the EHR (info card, suggestion, link)

Prior Authorization Workflow Simulation: Models payer-specific PA rules for selected services, simulating the electronic prior authorization (ePA) exchange using the Da Vinci Prior Authorization Support (PAS) FHIR Implementation Guide.

ONC 21st Century Cures Compliance Checklist: A 45-item checklist evaluating a hypothetical EHR system's compliance with ONC Health IT Certification criteria, including FHIR API requirements, information blocking prohibitions, and standardized API obligations.

Learning application:

The FHIR Lab is the bridge between policy (21st Century Cures requirements) and technology (actual FHIR implementation). An analyst who understands the regulatory requirements but cannot read a FHIR resource or understand a CDS Hooks exchange will struggle to evaluate whether a vendor is truly compliant.

Tool 2: Risk Stratification Engine

The Risk Stratification Engine applies the HCC v28 risk adjustment model to patient cohorts, enabling population segmentation by risk tier and comorbidity pattern analysis.

Core capabilities:

HCC v28 RAF Scoring: Enter a set of ICD-10 diagnoses for a patient; the engine maps them to HCC categories (respecting the hierarchical structure — only the most severe condition in each disease hierarchy contributes to the RAF score), applies age/sex coefficients, calculates interaction terms, and produces a composite RAF score.

The HCC hierarchy is conceptually important: a patient with both Type 2 Diabetes Without Complications (HCC 19) and Type 2 Diabetes With Chronic Complications (HCC 18) is scored only for HCC 18 — the more severe condition. This prevents double-counting and ensures the RAF score reflects the patient's most significant conditions.

Population Segmentation: Define a cohort (by uploading a condition list or using the built-in population simulator), and the engine segments the population into risk tiers:

Tier	RAF Range	Approximate % Population	Approximate % Cost
Complex	≥3.0	3%	40%
High Risk	1.5–2.9	12%	30%
Rising Risk	0.8–1.4	25%	20%
Low Risk	<0.8	60%	10%

Comorbidity Scoring: Two widely used comorbidity indices are implemented:

- **Charlson Comorbidity Index:** 17 conditions weighted by their association with 1-year mortality; score of 0–33
- **Elixhauser Comorbidity Set:** 30 conditions associated with in-hospital outcomes; used in hospital quality research

Both indices are useful for research, publication, and risk-adjustment in quality measurement contexts where the HCC model is not applicable.

Cohort Comparison: Enter two populations' diagnosis distributions and view side-by-side RAF score profiles, cost concentration curves, and high-condition prevalence comparisons.

Clinical application:

For a Vermont health system entering an MSSP contract, the Risk Stratification Engine enables:

1. Estimating the attributed population's RAF score distribution
2. Projecting the shared savings benchmark (dependent on population risk)
3. Identifying the high-complexity tier (3% of population, 40% of cost) for intensive care management targeting
4. Modeling the ROI of care management investment at different care management penetration rates

Section 2: Payment Models & VBC

Tool 3: APM Design Lab

The APM Design Lab allows analysts to design alternative payment models from first principles — not selecting an off-the-shelf CMMI model but constructing the model architecture appropriate for a specific organizational context.

Core capabilities:

APM Type Selection: Choose from:

- **Episode bundles (BPCI-A style):** Fixed payment for a defined clinical episode, including post-acute care
- **MSSP-style shared savings:** Benchmark-based shared savings with configurable MSR and sharing rates
- **ACO REACH global option:** Full prospective capitation for the attributed population
- **Primary care payment model:** Prospective enhanced primary care PMPM with quality performance adjustments
- **Global budget:** All-payer spending target for a defined geographic population

Benchmark Construction: For any APM type, configure:

- Historical trend rate for benchmark trend adjustment
- Risk adjustment methodology
- Quality withhold percentage and performance thresholds
- Risk corridor limits (cap on upside/downside)

Population Parameter Inputs: Enter the population characteristics that determine financial feasibility:

- Payer mix (Medicare, Medicaid, commercial %)
- Age/sex distribution
- Chronic condition prevalence (diabetes, CHF, COPD, behavioral health)
- Current utilization rates (admissions/1,000, ED visits/1,000)
- Current primary care-to-specialist visit ratio

Model Recommendation Engine: Based on population parameters, organizational characteristics (hospital size, current risk capability, data infrastructure), and market context, the engine produces a natural-language APM recommendation with financial rationale.

Vermont application:

For Vermont's community hospitals transitioning under AHEAD, the APM Design Lab can model the specific financial mechanics of their AHEAD budget allocation — what rate of cost reduction is required to achieve positive margin under their allocated budget, which clinical programs generate the most cost reduction per dollar invested, and what the financial impact of different benchmark methodologies would be.

Tool 4: APM Shared Savings Calculator

The APM Shared Savings Calculator implements the exact financial mechanics of MSSP tracks, ACO REACH options, and custom global budget scenarios.

MSSP Track Implementation:

MSSP TRACK OPTIONS			
Track	Sharing Rate	Risk Corridor	MSR
Basic A	40%	0%/0%	2–3.8%
Basic B	40%	0%/0%	2–3.8%
Basic C	50%	10%/10%	2–3.8%
Basic D	50%	10%/10%	2–3.8%
Basic E	50%	15%/15%	2–3.8%
Enhanced	75%	20%/20%	0% (no MSR)

Key calculation outputs:

- Gross savings vs. benchmark
- MSR check (does savings percentage exceed the minimum savings rate?)
- Quality performance score and multiplier
- Net shared savings payment
- Downside risk calculation (for two-sided risk tracks)
- Sensitivity analysis: medical cost trend ±2% (showing breakeven scenario)

Illustrative Vermont calculation:

VERMONT COMMUNITY HOSPITAL – MSSP ENHANCED TRACK SCENARIO	
Attributed population:	8,500 Medicare beneficiaries
RAF-adjusted benchmark PMPM:	\$876

Annual benchmark:	\$89.2M
Performance year actual spend:	\$84.8M
Gross savings:	\$4.4M (4.9% of benchmark)
No MSR for Enhanced Track:	Savings count from dollar 1
Sharing rate:	75%
Gross savings × rate:	\$3.3M
Quality multiplier:	×1.08 (strong quality performance)
Net shared savings:	\$3.56M
Net margin impact:	+1.8 percentage points
(Before sharing: −2.1% margin; After: −0.3% margin)	
This scenario shows that even a community hospital with a −2.1% operating margin can approach breakeven through MSSP shared savings under strong quality performance. The clinical prerequisite: reducing \$4.4M in spending vs. benchmark (primarily through reduced preventable admissions and post-acute care redesign).	

Tool 5: Cost-Effectiveness Analysis Calculator

The CEA Calculator implements the full methodology of cost-effectiveness analysis, from basic ICER calculation through probabilistic sensitivity analysis.

Core capabilities:

ICER Calculation:

- Input: incremental cost, incremental QALYs (or life years gained), comparator description
- Output: ICER (\$/QALY or \$/LY), threshold comparison, cost-effectiveness plane position

Number Needed to Treat (NNT):

- Input: event rate in intervention group, event rate in control group
- Output: NNT with 95% confidence interval, NNH (Number Needed to Harm)

Break-even Timeline:

- Input: upfront intervention cost, annual ongoing cost, annual avoided costs, discount rate
- Output: years to break even, NPV over 5 and 10 years

One-way sensitivity analysis: Tornado chart showing which input parameter (cost, effectiveness, discount rate, etc.) has the largest impact on the ICER.

Monte Carlo Probabilistic Sensitivity Analysis: 1,000 iterations sampling from cost and effectiveness probability distributions (log-normal for cost, beta for effectiveness proportions) producing:

- Cost-effectiveness acceptability curve (probability of cost-effectiveness at each willingness-to-pay threshold)
- Cost-effectiveness plane scatter plot
- EVPI (Expected Value of Perfect Information) — how much would it be worth to commission a definitive study before deciding?

Why EVPI matters:

EVPI quantifies the value of uncertainty resolution. If a decision involves a \$50M investment and the current evidence suggests the intervention is cost-effective with 65% probability, the EVPI tells you the maximum it would be worth paying for a definitive study that would tell you for certain. If EVPI is \$500K, you should commission a study; if EVPI is \$20K, the research investment is not justified — just make the decision.

Section 3: Population & Equity

Tool 6: Population Health Modeler

The Population Health Modeler implements three analytical frameworks for understanding disease burden and intervention impact in defined populations.

Framework 1: Markov Disease Progression Models

Markov models simulate the progression of chronic diseases through defined health states over time, enabling projection of future disease burden and intervention impact.

Five chronic disease models are implemented:

MARKOV MODEL STRUCTURE (Diabetes Example)

States:

[1] No diabetes

[2] Pre-diabetes (IFG/IGT)

[3] Diabetes, controlled (HbA1c ≤8%)

[4] Diabetes, uncontrolled (HbA1c >8%)

[5] Diabetes with complications

[6] Dead

Transition probabilities (annual, example values):

2→1: 0.15 (regression with lifestyle intervention)

2→3: 0.10 (progression to diabetes)

3→4: 0.25 (control deterioration)

4→3: 0.15 (control improvement with intervention)

4→5: 0.08 (complication development)

5→6: 0.05 (excess mortality from complications)

With care management intervention (assumptions):

2→1: 0.22 (+47% regression rate)

3→4: 0.18 (−28% control deterioration)

4→3: 0.22 (+47% control improvement)

4→5: 0.05 (−37% complication rate)

Users configure the model with their population's starting disease state distribution, select an intervention (or design a custom one with specified transition probability changes), and receive 5-year projections of:

- Projected disease state distribution at each time step
- Cumulative QALYs gained
- Cost trajectories (care costs by state × state population)
- Net present value of the intervention

Framework 2: SIR/SEIR Epidemic Models

The Susceptible-Infected-Recovered (SIR) model and its extension (SEIR, adding an Exposed state) are the foundational frameworks for epidemic dynamics. The tool allows:

- Configurable R_0 (basic reproduction number)
- Population size and initial infected count
- Vaccination rate (reducing susceptible population)
- Intervention effect (reducing transmission rate)
- Seasonal variation in transmission

Framework 3: Preventable Hospitalization Reduction Model

Input: current preventable hospitalization rates (PQIs), proposed care management program details, expected program penetration. Output: projected hospitalization reduction, QALY gain, and cost savings. Benchmarked against AHRQ PQI national and state rates.

Tool 7: Health Equity Studio

The Health Equity Studio consolidates five equity analytics into a single workflow, supporting the full arc from disparity identification through intervention prioritization.

Disparity Analysis:

Ten clinical and access outcomes are analyzed by race/ethnicity, income quartile, and rural/urban:

1. Diabetes control (HbA1c $\leq$ 8%)
2. Hypertension control (<140/90)
3. Colorectal cancer screening completion
4. Breast cancer screening completion
5. Childhood immunization series completion
6. Depression treatment receipt
7. Asthma control (ER visit rate)
8. Prenatal care (first trimester initiation)
9. Primary care access (annual visit)
10. ED utilization rate (preventable visits)

For each measure, the studio calculates:

- Absolute disparity (percentage point gap between highest and lowest-performing group)
- Relative disparity (ratio of lowest-to-highest group performance)
- HTR Disparity Priority Score (weighted by disparity magnitude, clinical significance, and intervention availability)

Geographic Access Mapping:

Drive-time analysis using HRSA HPSA data and geographic distance calculations:

- % of population within 30 minutes of nearest primary care physician
- % of population within 30 minutes of nearest emergency department
- % of population within 60 minutes of nearest OB/GYN
- Access gap map: geographic visualization of zip codes below access thresholds

SDOH Burden Score:

County-level composite of:

- Housing instability (% of households with severe housing cost burden)
- Food insecurity (% of households food insecure)
- Transportation barriers (% of households without vehicle access)
- Income poverty (% below 200% FPL)

HEROI Calculation:

As described in Chapter 8, the Health Equity Return on Investment metric applies an equity weighting to standard CEA analysis. The studio automates this calculation, allowing analysts to compare interventions on both standard ICER and equity-weighted HEROI.

Equity Action Plan Generator:

The studio's synthesis tool ranks potential equity interventions by:

1. Equity impact score (magnitude of disparity reduction $\times$ population affected)
2. Implementation feasibility (evidence base strength $\times$ local resource availability)
3. HEROI (equity-weighted cost-effectiveness)

Output: a prioritized action plan with specific program recommendations, evidence citations, and estimated implementation investment.

Section 4: Policy & Quality Sciences

Tool 8: Policy Simulator

Described in Chapter 4, the Policy Simulator models six policy scenario types: 1115 waiver design, global budget, Medicaid expansion, price transparency policy, prior authorization reform, and state comparison.

The global budget modeler in detail:

The Vermont global budget scenario is the most complex and most used Policy Simulator application. Users configure:

Population parameters:

- Covered population by payer (Medicaid, Medicare Advantage, commercial)
- FFS Medicare inclusion (on/off — illustrating the budget gap created by exclusion)
- Age and chronic disease distribution

Budget parameters:

- Base year per-capita PMPM by payer
- CMMI-negotiated trend rate (growth cap)
- Quality risk corridor (bonus/penalty for meeting/missing quality targets)
- Hospital-specific budget allocation methodology (equal per-capita, need-adjusted, historical-based)

Output:

- 5-year budget trajectory by payer
- Per-hospital budget allocation and financial impact
- Required cost reduction per year to stay within budget
- Equity impact analysis (differential budget pressure by hospital type)
- Break-even analysis (at what cost trend does the system stay within budget?)

Tool 9: Clinical Quality Optimizer

The Clinical Quality Optimizer translates raw quality measure data into strategic quality improvement priorities and financial projections.

HEDIS Simulation:

Enter current performance rates for 15 HEDIS measures; the tool benchmarks each against NCQA's national percentile distribution and calculates:

- Current percentile ranking
- Gap to NCQA 75th percentile (the standard quality threshold in most VBC contracts)
- Gap to NCQA 90th percentile (advanced quality threshold in some contracts)
- Point impact per 1% improvement in measure rate

CMS Stars Predictor:

Enter current performance rates for CMS Star measure sub-components; the tool projects the overall Star Rating using CMS's star cut-point methodology. Critically, it identifies which measures are closest to the next cut-point — where marginal improvement effort yields a Star Rating jump.

STAR RATING SENSITIVITY ANALYSIS EXAMPLE

Current projected rating: 3.5 stars

Measures closest to next cut-point:

Annual Flu Vaccine:

0.3% from 4-star cutpoint

Medication Adherence-Statins:

0.8% from 4-star cutpoint

HbA1c Testing (diabetes):

0.1% from 4-star cutpoint

If all three improved to 4-star cutpoint:
Projected new rating: 4.0 stars
Financial impact at 100,000 enrollees:
Quality Bonus Payment (+5%): +\$63M/year in CMS revenue

MIPS Optimizer:

MIPS composite score optimization across 4 categories (Quality, Cost, Improvement Activities, Promoting Interoperability) by:

- Selecting quality measures that maximize points given the practice's patient population
- Identifying improvement activities that qualify for maximum credit with minimum operational burden
- Calculating the Expected Payment Adjustment (bonus or penalty) from the projected MIPS score

P4P Revenue Calculator:

Enter the structure of a commercial payer's Pay-for-Performance bonus program (target measures, bonus tiers, base revenue at risk), enter current and projected performance rates, and calculate the P4P revenue impact of reaching each performance tier.

Tool 10: Hospital Financial Scorecard

The Hospital Financial Scorecard is a comprehensive financial stress-testing and benchmarking tool for hospital executives and analysts.

Payer Mix Analysis:

HOSPITAL PAYER MIX ANALYZER

Input your hospital's payer mix (% of revenue by payer) and each payer's current rate as % of Medicare:

	Your Hospital	CAH Median	Rural PPS	Urban Tertiary
Medicare	42%	48%	38%	30%
Medicaid	24%	22%	19%	16%
Commercial	29%	27%	38%	50%
Self-pay	5%	3%	5%	4%

Medicare rate:

100% of Medicare

→ benchmark

Medicaid rate:

85% of Medicare

→ loss-maker

Commercial rate:

285% of Medicare

→ cross-subsidy

Net revenue per adjusted discharge:
Your Hospital: \$8,420 (vs. Operating Cost: \$8,912)
Operating Margin: -5.5%

Medicaid Rate Sensitivity:

Slider from -10% to +10% Medicaid rate change → net operating revenue impact. At Vermont's current payer mix, a 10% Medicaid rate reduction reduces system operating revenue by approximately \$8.3M — enough to push most community hospitals below -5% operating margin.

Volume Stress Test:

Model inpatient and outpatient volume changes (from -30% to +30%) and their impact on operating margin. Critical insight: for high-fixed-cost hospitals, volume reductions have disproportionately large margin impacts because fixed costs do not decline proportionally with volume.

Peer Benchmarking:

Compare your hospital's financial metrics against three peer groups:

- Critical Access Hospital (CAH): small, rural, cost-based Medicare payment
- Rural PPS: small-medium rural hospital under prospective payment
- Urban Tertiary: large, urban academic or community hospital

The benchmarking reveals where your hospital's performance is consistent with peers (operating in a challenging market) versus where it is an outlier (potential for operational improvement).

DSH and 340B Analysis:

- Disproportionate Share Hospital (DSH) payment estimate based on your Medicaid/uninsured payer mix
- 340B drug discount program savings estimate (applicable if your hospital qualifies for 340B status based on DSH percentage or other qualifying criteria)

Tool 11: HTA Studio

The Health Technology Assessment Studio is the most methodologically sophisticated tool in the Research Lab, implementing the full HTA methodological framework used by ICER, NICE, and CDA-AMC.

Budget Impact Model:

Projects the 5-year financial impact of a new technology (drug, device, diagnostic, or care model) on a payer or health system's budget:

- Incidence of target condition in the payer's population
- Market uptake trajectory (S-curve adoption)
- Technology cost (net of any offsets from displaced treatments)
- Safety/adverse event costs
- Net annual budget impact per year 1–5

Multi-Criteria Decision Analysis (MCDA):

When cost-effectiveness alone is insufficient for decision-making (e.g., when a high-cost therapy has transformative benefits for a small population), MCDA provides a structured multi-criteria framework:

Eight criteria are evaluated:

1. Clinical effectiveness (ICER vs. threshold)
2. Safety profile (adverse event rate)
3. Health equity impact (differential benefit for disadvantaged populations)
4. Burden of illness (severity and unmet need of the condition)
5. Evidence quality (strength and generalizability of the evidence base)
6. Innovation (mechanistic novelty or paradigm shift)
7. Patient/caregiver value (impact on quality of life and caregiver burden)
8. Cost to payer (total budget impact)

Each criterion is scored 0–10; criteria are weighted by the decision-maker (weights must sum to 100%); the weighted composite score provides a structured basis for coverage or formulary decisions.

Monte Carlo PSA:

1,000 iterations sample from parameter probability distributions:

- Cost: log-normal distribution (log-mean, log-SD)
- Effectiveness (QALYs): beta distribution (alpha, beta parameters)
- Other parameters: user-specified distributions

Output: cost-effectiveness plane scatter plot with quadrant distribution, cost-effectiveness acceptability curve, and EVPI calculation.

Tool 12: Actuarial Lab

The Actuarial Lab covers the core methodological tools of health insurance actuarial analysis.

ACA Actuarial Value Calculator:

Enter a plan's benefit design (deductible, copays, coinsurance, OOP maximum, covered services) and the calculator estimates the plan's actuarial value using CMS's actuarial value calculator methodology. Used to verify metal tier classification and to compare benefit designs.

Premium Development:

Three rating methodologies are implemented:

- 1. **Community rating:** Single premium for all enrollees (adjusted only for geography and tobacco status)
- 2. **Age-banded rating:** Premium varies by age using the ACA 3:1 age band limit
- 3. **Experience rating:** Premium based on the group's own historical claims experience (primarily for large self-insured employers)

Adverse Selection Model:

Simulate enrollment dynamics under different premium levels and benefit designs. Adjust health status distribution assumptions (what percentage of enrollees are healthy vs. sick), and see how the risk pool changes year-over-year — including the trajectory toward or away from death spiral conditions.

IRA 2022 Drug Pricing Impact:

Model the estimated negotiated drug price impact under the Inflation Reduction Act's Medicare drug price negotiation provisions, by drug category (cardiovascular, diabetes, oncology, specialty) and payer type (Medicare, commercial, Medicaid).

Medical Loss Ratio Analysis:

Calculate the MLR for a health plan (medical expenses as % of premium revenue), determine whether rebates are owed to members (if below 80% small group or 85% large group), and model the premium adjustment required to achieve target MLR.

Section 5: Technology & AI

Tool 13: AI Analytics Lab

The AI Analytics Lab evaluates clinical AI tools using standardized methodologies from the FDA, NCQA, and emerging AI governance frameworks.

Model Performance Comparison:

Compare up to four predictive models side-by-side on:

- AUC-ROC (Area Under the Receiver Operating Characteristic Curve): ability to discriminate positive from negative cases
- Sensitivity (recall): true positive rate
- Specificity: true negative rate
- Precision (positive predictive value): proportion of positive predictions that are correct
- F1 score: harmonic mean of precision and recall

MODEL COMPARISON EXAMPLE – SEPSIS PREDICTION				
	Model A (EHR-ML)	Model B (qSOFA)	Model C (NEWS2)	Model D (SOFA)
AUC-ROC	0.87	0.73	0.79	0.74
Sensitivity	0.82	0.65	0.74	0.71
Specificity	0.79	0.82	0.76	0.77
Precision	0.41	0.38	0.44	0.40
F1 Score	0.55	0.49	0.57	0.52

Bias Analysis (Black vs. White patients):

- AUC-ROC gap: 0.09 (Model A underperforms for Black patients)
- EO Diff: 0.11 (sensitivity gap – Black patients less likely to be flagged as high-risk)

Action required: Before clinical deployment, Model A needs bias mitigation and targeted validation on Black patient cohort.

Bias Detection:

Three bias metrics (described in Chapter 6) are calculated for each model against user-specified demographic groups.

AI Governance Checklist:

A 40-item governance framework covering:

- Data governance (collection methods, consent, representativeness)
- Model development (validation methodology, external validation, performance documentation)
- Clinical integration (workflow integration, override capability, alert fatigue mitigation)
- Ongoing monitoring (performance drift detection, disparate impact monitoring)
- Explainability (feature importance documentation, patient/clinician communication)

ROI Calculator:

For an AI implementation investment:

- Input: implementation cost (license + integration + training), productivity gains (FTE-hours saved/year), error reduction savings (avoided adverse events × cost per event)
- Output: net present value over 5 years, payback period, ROI %

Tool 14: Digital Health Lab

RPM ROI Calculator:

Remote Patient Monitoring (RPM) is a growing modality for managing chronic disease between office visits. CMS reimburses RPM under CPT codes 99453–99458:

RPM CPT CODE	REFERENCE	
99453	Initial setup and patient education:	\$19
99454	Device supply, 30 days transmission:	\$48/month
99457	Remote monitoring, 20+ min/month:	\$51/month
99458	Additional 20 min monitoring:	\$42/month
Average monthly RPM revenue per patient:		\$141
(At 60 minutes of remote monitoring per month)		
RPM ROI Calculator: enter panel size, device cost, monitoring staff hours, readmission reduction rate → net annual revenue and cost savings		

For a Vermont community hospital managing 500 CHF patients with RPM:

- Annual RPM revenue: $500 \times \$141 \times 12 = \$846,000$
- Device and staff cost: \$380,000
- Net margin: \$466,000
- Readmission reduction (15% reduction in CHF readmissions): \$225,000 in avoided costs
- Total annual benefit: \$691,000

Telehealth Modeling:

Five CMS policy scenarios model the impact on telehealth utilization and revenue:

1. Geographic restrictions restored (originating site requirement)
2. Current flexibilities made permanent
3. Audio-only telehealth covered at parity with video
4. Mental health telehealth carved out with separate policy
5. FQHCs and RHCs included in permanent telehealth payment

Patient Engagement Platform Scorecard:

A 12-criterion scorecard comparing digital front door, patient portal, and remote monitoring vendors across: clinical integration (FHIR, EHR), patient activation features, language access, accessibility (ADA), analytics capability, security/HIPAA compliance, mobile experience, and vendor financial stability.

Section 6: Knowledge & Workspace

Tool 15: Evidence Library

The Evidence Library is the Research Lab's reference database — a curated, searchable collection of landmark research organized into three collections:

CEA/CUA Study Database (25 studies):

Landmark cost-effectiveness analyses in healthcare transformation domains, filterable by:

- Condition (diabetes, cardiovascular, behavioral health, cancer, etc.)
- Intervention type (medication, care model, technology, prevention)
- ICER threshold comparison
- Evidence strength (RCT, retrospective cohort, modeling)

Each study entry includes:

- Citation (AMA and APA format)
- Population and setting
- Comparator
- ICER (\$/QALY)
- Limitations
- HTR relevance note

CMMI Model Tracker (20 models):

Evaluation summaries for 20 CMMI models, including:

- Model description and program structure
- Current status (active, concluded, suspended, expanded)
- Evaluation findings (savings generated, quality impacts, participation)
- Lesson learned summary
- Applicability to Vermont context

HTR Policy Brief Library (15 briefs):

Platform-authored policy analysis briefs including:

- Structured findings and evidence summary
 - Recommendations (with evidence grading)
 - Vermont applicability assessment
 - Related Research Lab tools
-

Tool 16: Workforce Modeler

Physician Pipeline Projector:

Project physician supply and demand across 12 specialties over 10 years using configurable parameters:

- Residency graduation rates (by specialty)
- Retirement curves (age-specific exit rates)
- International medical graduate (IMG) contribution
- Demand growth rate (demographic + utilization trend)

Output: 10-year supply/demand gap by specialty, with and without intervention scenarios (loan repayment, scope-of-practice expansion, APRN substitution).

Nurse Staffing Ratio Impact Model:

Models patient outcomes and cost impact under mandatory nurse staffing ratio legislation, drawing on California's mandate (1:5 RN:patient in medical-surgical) as the policy benchmark. Calculates:

- Additional RN FTEs required to meet proposed ratio
- Annual cost of compliance
- Evidence-based outcome improvement (readmission reduction, mortality reduction)
- Net cost-benefit (outcome improvement savings vs. staffing cost)

Turnover Cost Calculator:

NURSING TURNOVER COST BREAKDOWN (per RN)	
Recruitment costs: (Job board, agency, recruiter fee)	\$8,200
Onboarding and orientation: (Training staff time, materials, new hire productivity ramp)	\$9,400
Vacancy coverage: (Agency temp, overtime for staff)	\$16,800
Lost productivity during ramp: (New hire at 60–70% productivity for first 3 months)	\$12,000
Total per-nurse turnover cost:	\$46,400 median (Range: \$28,000–\$88,000 depending on specialty and market)
At 22% turnover rate in 500-nurse hospital:	
Annual turnover: 110 nurses	
Annual cost: 110 × \$46,400 = \$5.1M	
Reducing to 15% turnover:	
Prevented turnovers: 35	
Cost savings: 35 × \$46,400 = \$1.6M/year	

Tool 17: Innovation Leaderboard

State Leaderboard:

All 50 states ranked on a composite health transformation index (policy innovation + VBC penetration + equity progress + technology adoption), with year-over-year change indicators and detailed dimension breakdowns. Supports custom weighting to prioritize the dimensions most relevant to a user's analytical focus.

Health System VBC Maturity:

30 major health systems rated across 6 VBC maturity dimensions:

1. Risk contract penetration (% of revenue in risk arrangements)

- 2. Care management program maturity (tier 1–4 complexity management)
- 3. Data infrastructure (analytics capability for population health)
- 4. Quality performance (HEDIS/Stars composite vs. peers)
- 5. Consumer experience (NPS, patient satisfaction scores)
- 6. Network design (specialist alignment, post-acute integration)

Payer Innovation Scorecard:

20 major payers rated on innovation leadership:

- APM adoption rate (% of contracts with value-based payment)
- Digital health investment (annual spend, programs offered)
- SDOH programs (community health worker, food insecurity, housing)
- Equity reporting (public disclosure of race/ethnicity stratified quality data)
- Interoperability readiness (FHIR API implementation maturity)

Tool 18: Research Workspace

The Research Workspace is the Research Lab's productivity layer — enabling analysts to save, organize, annotate, and export their work.

Scenario Manager: Save named parameter sets from any Research Lab tool. When you return to the APM Shared Savings Calculator with a different population profile next month, you can load the previous scenario for comparison rather than re-entering all parameters.

Report Builder: Structured templates for common deliverable types:

- Executive Summary (audience, question, findings, recommendations)
- Methodology Note (data sources, analytical approach, limitations)
- Policy Brief (context, evidence, recommendations, implications)
- Business Case (problem statement, proposed solution, financial model, risks)

Citation Manager: Maintains a bibliography in AMA and APA format for all sources used across Research Lab tools and manual additions. Generates formatted bibliography for inclusion in reports.

Export: Markdown (.md) or plain text (.txt) download of any report or scenario summary.

Key Concepts Introduced in This Chapter

- Markov disease progression model:** A mathematical model simulating patient transitions between defined health states over time, used to project disease burden and the impact of clinical interventions.
- EVPI (Expected Value of Perfect Information):** The maximum value of commissioning further research before making a decision under uncertainty; calculated from probabilistic sensitivity analysis.
- MCDA (Multi-Criteria Decision Analysis):** A decision framework evaluating options against multiple weighted criteria simultaneously; used in HTA when cost-effectiveness alone is insufficient.
- RPM (Remote Patient Monitoring):** Delivery of healthcare monitoring services via connected devices in the patient's home; reimbursed by CMS under CPT codes 99453–99458.
- AUC-ROC:** Area Under the Receiver Operating Characteristic Curve; a measure of a predictive model's ability to discriminate between positive and negative cases across all classification thresholds.
- Gold-carding:** A practice or policy exempting providers with consistently high prior authorization approval rates from individual PA review for specific services.

CHAPTER 14: Learning for Transformation — The HTR Academy

Why Learning Infrastructure Matters

Healthcare transformation fails when knowledge fails. Not when the policy is wrong or the payment model is incorrectly structured — although those failures happen too — but when the people responsible for implementing transformation do not understand the concepts, frameworks, and evidence that should guide their decisions.

A hospital CFO who does not understand HCC risk adjustment will make revenue cycle decisions that leave millions on the table under a Medicare Advantage contract. A quality improvement manager who does not understand HEDIS denominator specifications will report the wrong patients in the numerator and denominator, producing invalid quality scores. A state Medicaid director who does not understand global budget benchmark methodology will negotiate a waiver that locks in the system's historical inefficiencies.

The HTR Academy exists to close these knowledge gaps systematically — at every professional level, in every pillar domain, through a learning architecture that is structured enough to build genuine expertise and flexible enough to serve diverse learners with different starting points.

The Academy Content Architecture

ACADEMY CONTENT HIERARCHY

TRACK

A curated learning path (4–12 hours total) for a defined audience (e.g., "Hospital CFOs entering first VBC contract")

COURSE

A structured multi-module program with assessment and certificate (typically 4–8 hours)

MODULE

The atomic learning unit (30–90 minutes)

Learning objectives (3–5)

Body content (rich text, video, audio)

Knowledge checks (3–5 questions)

Reflection prompts

Supplementary content:

WEBINARS – Live and on-demand expert sessions

CASE STUDIES – Real-world implementation analyses

GLOSSARY – 200+ term definitions, pillar-tagged

PERSONALIZED LEARNING – AI-generated individual curricula

The Six Learning Tracks

Track 1: Policy Analyst Track

Audience: Policy analysts, government affairs professionals, Medicaid agency staff **Duration:** ~10 hours **Core competencies:**

- ACA architecture and insurance market mechanics
- CMMI model portfolio and evaluation methodology
- Medicaid waiver design (1115, 1332, 1115A)
- State vs. federal policy levers
- Vermont Act 167 and AHEAD Model policy analysis

Sample module sequence:

1. *Understanding the ACA's Structural Architecture* (Foundational, 45 min)
2. *CMMI and the Federal Payment Reform Laboratory* (Intermediate, 60 min)
3. *Medicaid Waiver Design: 1115 and Beyond* (Intermediate, 75 min)
4. *Vermont's Policy Transformation: Act 167 in Context* (Advanced, 60 min)

5. *Simulation: Using the Policy Simulator for Waiver Impact Analysis* (Applied, 45 min)

Track 2: Health Economist Track

Audience: Health economists, financial analysts, actuaries, managed care executives **Duration:** ~12 hours **Core competencies:**

- VBC economics and APM financial mechanics
- Cost-effectiveness analysis methodology
- Actuarial principles for health insurance
- Hospital financial analysis and stress-testing
- Global budget economics (Vermont AHEAD model)

Sample module sequence:

1. *The Economics of Value-Based Care* (Foundational, 60 min)
 2. *APM Financial Mechanics: MSSP to ACO REACH* (Intermediate, 90 min)
 3. *Cost-Effectiveness Analysis: ICER, QALY, and Threshold Analysis* (Advanced, 75 min)
 4. *Actuarial Principles in Health Insurance* (Intermediate, 60 min)
 5. *Hospital Financial Analysis: Reading the Balance Sheet* (Intermediate, 60 min)
 6. *Vermont Global Budget Economics* (Advanced, 60 min)
-

Track 3: Clinical Leader Track

Audience: CMOs, medical directors, quality improvement officers, clinical informaticists **Duration:** ~10 hours **Core competencies:**

- HEDIS, CMS Stars, and MIPS quality measurement frameworks
- Risk stratification and HCC coding
- Care delivery model redesign (PCMH, ACO)
- Workforce planning and scope of practice
- Behavioral health integration

Sample module sequence:

1. *Quality Measurement Frameworks: HEDIS, Stars, MIPS* (Foundational, 75 min)
 2. *Risk Stratification: HCC v28 and Population Segmentation* (Intermediate, 60 min)
 3. *Redesigning Primary Care: PCMH and ACO Models* (Intermediate, 60 min)
 4. *Behavioral Health Integration: The Collaborative Care Model* (Advanced, 75 min)
 5. *Workforce Planning Under Healthcare Transformation* (Intermediate, 60 min)
-

Track 4: Technology Strategist Track

Audience: CIOs, health IT directors, digital health product managers, interoperability engineers **Duration:** ~12 hours **Core competencies:**

- FHIR R4 architecture and US Core Implementation Guide
- 21st Century Cures Act and information blocking
- Clinical AI governance and bias assessment
- Telehealth policy and technology infrastructure
- EHR strategy and vendor evaluation

Sample module sequence:

1. *FHIR R4: From Standard to Implementation* (Intermediate, 90 min)
 2. *21st Century Cures and Information Blocking* (Foundational, 45 min)
 3. *Clinical AI: Governance, Bias, and Validation* (Advanced, 75 min)
 4. *Telehealth Technology and Policy* (Intermediate, 60 min)
 5. *Building a FHIR Resource — Hands-On Lab* (Applied, 60 min)
-

Track 5: Health Equity Track

Audience: Equity officers, community health workers, SDOH program managers, population health analysts

Duration: ~10 hours **Core competencies:**

- Health disparities measurement and interpretation
- SDOH frameworks and intervention design
- Equity-weighted cost-effectiveness analysis (HEROI)
- NCQA Health Equity Accreditation standards
- Community health worker program design

Sample module sequence:

1. *Understanding Health Disparities: Evidence and Mechanisms* (Foundational, 60 min)
 2. *Social Determinants of Health: Framework and Intervention* (Foundational, 75 min)
 3. *Measuring Health Equity: NCQA HES, CMS HEI, HEROI* (Intermediate, 60 min)
 4. *SDOH Screening in Clinical Settings* (Applied, 45 min)
 5. *Community Health Worker Program Design* (Advanced, 60 min)
 6. *Vermont Equity Case Study: Rural Health and LGBTQ+ Access* (Applied, 45 min)
-

Track 6: Operations Leader Track

Audience: COOs, revenue cycle directors, supply chain managers, compliance officers **Duration:** ~10 hours **Core competencies:**

- Revenue cycle fundamentals and VBC-era RCM
- HCC coding optimization and RAF gap management
- Prior authorization workflow and electronic PA
- Workforce credentialing and operational stability
- HIPAA compliance and healthcare operations security

Sample module sequence:

1. *Revenue Cycle Management: FFS to Value-Based* (Foundational, 75 min)
 2. *HCC Coding and RAF Optimization* (Intermediate, 60 min)
 3. *Prior Authorization: Operations, Compliance, and Reform* (Intermediate, 60 min)
 4. *Workforce Operations: Credentialing and Retention* (Intermediate, 60 min)
 5. *VBC Operational Readiness Assessment* (Applied, 60 min)
-

Case Studies: Learning from Real-World Implementation

Case studies are the most distinctive content type in the Academy — real-world analyses of healthcare transformation implementations, structured to extract generalizable lessons rather than simply tell success stories.

Case Study Structure:

Each case study follows a standardized structure:

1. **Context:** Organization background, market environment, what problem they were trying to solve
2. **Approach:** What they did, key decisions, timeline, partners
3. **Results:** Quantified outcomes with baselines and timeframes (using structured key outcomes cards)
4. **Lessons Learned:** What worked, what failed, what they would do differently
5. **Applicability:** Who can use this approach and under what conditions

Sample Case Study: Vermont FQHC Behavioral Health Integration

Organization: Multi-site Federally Qualified Health Center (FQHC) in rural Vermont (5 sites, 12,000 patients)

Context: The FQHC served a predominantly low-income, uninsured, and Medicaid population in three rural counties. Mental health comorbidities in 43% of patients were driving high ED utilization and poor chronic disease

control. A 6-month wait for specialty mental health referrals meant most patients never received treatment.

Approach: Implemented the Collaborative Care Model (CoCM) across all 5 sites over 18 months:

- Hired a part-time psychiatric consultant for each site (1 psychiatrist covering 5 sites via scheduled telephone consultation)
- Hired 5 behavioral health care managers (BHCMS) — one per site — with master's-level training
- Integrated PHQ-9 screening and GAD-7 into annual wellness visits
- Created a caseload registry tracking treatment initiation, response (PHQ-9 change), and remission

Key Outcomes:

Metric	Baseline	12-Month Result
Depression treatment initiation rate	23%	68%
PHQ-9 ≥5 point improvement	31%	59%
ED visits for psychiatric crisis	8.3/100 patients/year	4.1/100
HbA1c control (diabetics with depression)	41%	57%

Lessons Learned:

- BHCM caseload must stay below 60 active patients; above that, engagement quality deteriorates
- PHQ-9 follow-up scheduling requires proactive outreach — patients will not self-schedule
- Billing workflow for CoCM CPT codes requires specific training; initially left significant revenue uncaptured
- Psychiatric consultant availability at specific times (twice weekly rounds) was the highest-value element

Applicability:

- Directly applicable to Vermont's Act 167-funded behavioral health integration program
- FQHC context is generalizable to any primary care practice serving high-SDOH populations
- The part-time psychiatric consultant model (one per 5 sites) is a cost-effective staffing model for rural settings

The Glossary: 200+ Terms for a New Vocabulary

The HTR Academy Glossary provides definitions for 200+ healthcare transformation terms — organized by pillar, cross-referenced to related terms, and linked to the modules and articles that use them. A sample across pillars:

Policy terms:

- *APM (Alternative Payment Model):* Any payment arrangement that ties reimbursement to quality, outcomes, or total cost of care rather than fee-for-service volume
- *1115 Waiver:* CMS authority allowing states to waive standard Medicaid requirements to test new approaches, subject to federal budget neutrality
- *CMMI:* Center for Medicare and Medicaid Innovation; the federal laboratory for testing new payment and care delivery models

Economics terms:

- *PMPM (Per Member Per Month):* The per-member cost or payment in a capitated or prospective payment arrangement
- *Benchmark:* The spending target against which actual spending is measured in a shared savings or global budget arrangement
- *Risk corridor:* The cap on upside and downside financial exposure in a VBC contract

Technology terms:

- *FHIR Resource:* A discrete, modular data object in the FHIR standard representing a specific clinical or administrative concept

- *EHR (Electronic Health Record)*: A digital version of a patient's medical history maintained across clinical encounters; the primary clinical data system in modern healthcare
- *Interoperability*: The ability of different information technology systems to exchange data and use the information exchanged

Clinical terms:

- *HEDIS*: Healthcare Effectiveness Data and Information Set; NCQA's set of 90+ measures used to evaluate health plan and provider performance
- *RAF (Risk Adjustment Factor)*: The HCC-derived score predicting a patient's expected healthcare costs relative to the average Medicare beneficiary
- *PQI (Prevention Quality Indicator)*: AHRQ's measures identifying hospitalizations preventable with adequate outpatient care

Equity terms:

- *SDOH (Social Determinants of Health)*: The conditions in which people are born, grow, live, work, and age that produce health disparities
- *HEROI (Health Equity Return on Investment)*: An equity-weighted cost-effectiveness metric that values interventions targeting disadvantaged populations more highly
- *CHW (Community Health Worker)*: Trained community member providing health education, navigation, and social needs assistance

Operations terms:

- *Clean claim rate*: The percentage of submitted claims accepted without error on first submission
- *Days in AR*: The average number of days between service delivery and payment collection
- *RAF coding gap*: The difference between a patient population's clinically justified RAF score and their submitted (coded) RAF score

Personalized Learning: AI-Generated Curricula

The Academy's Personalized Learning feature represents the convergence of the HTR platform's AI capabilities with its educational content. Rather than selecting from a menu of pre-built tracks, a learner describes their role, knowledge gaps, and learning goals, and the AI generates a custom multi-week curriculum drawing from across the full Academy content library.

The onboarding form captures:

1. **Role:** Hospital Administrator, Policy Analyst, Clinician, Health Economist, Operations Manager, State Agency Staff, Health Technology Professional, Researcher/Academic, Other
2. **Topics of Interest:** Multi-select from all pillar areas and sub-topics
3. **Difficulty level:** Foundational / Intermediate / Advanced
4. **Available hours per week:** 2–4 / 5–8 / 9–12 / 12+
5. **Learning goals:** Free text (up to 300 words) — what you want to be able to do or understand after completing the curriculum
6. **Format preferences:** Readings, case studies, knowledge checks, reflections

The AI curriculum generation process:

The onboarding form data is sent to the backend with a structured prompt to the AI:

"You are an expert curriculum designer for healthcare transformation education. Based on the following learner profile, generate a detailed 8-week learning curriculum... The curriculum must draw from the following available content library items... Each week should have a theme, 2–4 specific learning items with estimated time, reflection questions, and a brief explanation of how the week's content connects to the learner's stated goals..."

The AI returns a structured JSON curriculum that is stored in the `user_learning_paths` database table and rendered as a week-by-week plan in the Academy interface.

Sample personalized curriculum output (abbreviated):

Learner profile: Hospital CFO, 20 years hospital finance experience, new to value-based care, needs to understand ACO financial mechanics and workforce operations for a planned MSSP Enhanced track transition.

PERSONALIZED CURRICULUM – WEEKS 1–3 (SAMPLE)

WEEK 1: From Fee-for-Service to Value-Based Payment

Theme: Understanding the economic architecture shift

Monday:

Reading: "The VBC Spectrum" (45 min)

→ Relevance: Foundation for understanding your MSSP contract

Wednesday:

Case Study: "Regional Health System – MSSP Year 1 Experience" (30 min)

Knowledge Check: 5 questions on MSSP financial mechanics

Friday:

Tool Lab: APM Shared Savings Calculator (45 min)

→ Enter your hospital's estimated attributed population and medical cost trend; see projected shared savings scenarios

Reflection: "What financial risks in the MSSP Enhanced track most concern your organization, and why?"

WEEK 2: HCC Coding and RAF Optimization

Theme: Revenue integrity under risk contracts

...

WEEK 3: Workforce Operations for VBC

Theme: Building the care coordination capacity VBC requires

...

Key Concepts Introduced in This Chapter

- HEDIS denominator:**

The eligible population for a HEDIS measure — correctly identifying who is in the denominator (eligible for the measure) is as critical as measuring performance for those patients.
- BHCM (Behavioral Health Care Manager):**

A master's-level clinician embedded in primary care under the Collaborative Care Model; manages a caseload of patients with depression and anxiety, tracks treatment response, and coordinates with the psychiatric consultant.
- CoCM (Collaborative Care Model):**

Evidence-based integrated behavioral health model reimbursed under CMS CPT codes 99492–99494; embeds psychiatric consultation and behavioral health care management in primary care.
- Personalized learning:**

AI-generated individual curricula based on a learner's role, goals, and knowledge gaps; draws from the full HTR Academy content library.
- Knowledge check:**

A brief embedded assessment within an Academy module, testing comprehension (not memorization) with explanation of the correct answer.

CHAPTER 15: AI and Healthcare Intelligence — The RAG-Powered Analyst

The Problem AI Solves in Healthcare Intelligence


Healthcare transformation generates an enormous volume of information: federal rulemaking, CMMI model evaluations, state policy changes, HEDIS updates, FHIR specification revisions, peer-reviewed research, and advisory analyses. A healthcare professional who needs to understand the implications of a pending CMS rule change, the current evidence on behavioral health integration, and Vermont's AHEAD model benchmark methodology — all in the same afternoon — faces an information retrieval challenge that no conventional search engine is designed to solve.

The HTR AI Analyst is designed to answer complex, multi-dimensional healthcare transformation questions by:

1. Searching a curated knowledge base of HTR analyses, policy documents, research, and Vermont-specific content
2. Retrieving the most relevant context for the specific question
3. Generating a grounded, cited response that is accurate to the retrieved evidence

This is **Retrieval-Augmented Generation (RAG)** — the combination of information retrieval and language model generation that produces answers that are both fluent (generated by a language model) and accurate (grounded in retrieved evidence rather than model memorization).

RAG Architecture: How the System Works



STREAMING RESPONSE → BROWSER

Tokens stream token-by-token; citations inline

The Knowledge Base: What the AI Knows

The AI Analyst's knowledge base is indexed from five source types:

1. HTR Policy Analyses (Sanity CMS)

All published HTR articles across the six pillar categories — deep-dive policy analyses, healthcare economics reviews, technology assessments, clinical quality analyses, equity research, and operations studies. These are the primary editorial content of the HTR platform, written by domain experts and updated continuously.

2. HTR Academy Content (Sanity CMS)

All published Academy modules, case studies, and glossary definitions — providing both conceptual educational content and applied case study analyses that the AI can draw on when answering practical "how" questions.

3. General Reference PDFs

Key reference documents indexed from the `/backend/data/` directory:

- *Wyman Report (Oliver Wyman/Marsh McLennan, 2024)*: Vermont Act 167 Community Engagement analysis — the most-referenced external document in the Vermont content
- *Health Economics reference*: Foundational health economics methodology text
- *AI Doctor*: Medical AI applications reference
- *Dictionary of Health Economics*: Terminology reference for economics queries

4. Vermont Medicaid Eligibility Documents

A specialized subsystem covers Vermont Medicaid eligibility — 23 official state documents covering MABD (Medicaid for Adults with Blindness or Disability), Dr. Dynasaur (children's Medicaid), Reach Up (cash assistance), Expanded Medicaid, and CHIP eligibility rules.

This subsystem uses a different chunking strategy than the general knowledge base: regulatory documents with conditional logic (IF age ≥ 65 AND income ≤ X THEN eligible) must be kept intact, not split at sentence boundaries, to preserve the eligibility reasoning structure.

5. HTR Methodology Documents

The HTI Methodology, White Paper, Research Lab documentation, and Advisory Services documentation are indexed to support questions about HTR's analytical frameworks and the platform itself.

The Medicaid Eligibility Subsystem

The Vermont Medicaid eligibility subsystem illustrates how a RAG system can be specialized for a high-stakes, rule-intensive domain.

The challenge:

Medicaid eligibility determination involves complex conditional logic: income thresholds, asset limits, categorical eligibility categories, immigration status requirements, age-specific rules, and program-specific exceptions. A language model that generates an eligibility answer from memorized training data — rather than from current, official eligibility rules — will produce unreliable results that could harm users who rely on them.

The solution:

The Medicaid eligibility path uses:

1. **Intent detection:** The system detects whether a query involves Medicaid eligibility keywords (income limits, MABD, Dr. Dynasaur, coverage eligibility, etc.) and routes it to the Medicaid path

- 2. **Scoped retrieval:** The search is filtered to the **Medicaid Eligibility** pillar only — retrieving from the 23 official Vermont Medicaid documents exclusively
- 3. **Expanded retrieval:** 8 top chunks (vs. 5 for general queries) are retrieved, because eligibility determinations require multiple rule sections to reason correctly
- 4. **Specialized system prompt:** The Medicaid system prompt requires step-by-step eligibility reasoning, mandatory rule citation (which document, which section), and a mandatory disclaimer directing the user to the Vermont Health Connect (VHC) for official determination

VERMONT MEDICAID ELIGIBILITY – SAMPLE QUERY

User: "I'm 38 years old, single, living in Vermont, with a monthly income of \$1,800. Am I eligible for Medicaid?"

AI Response:

Based on the Vermont Medicaid eligibility rules, here is an analysis of your potential eligibility:

EXPANDED MEDICAID (ACA Expansion):
Vermont expanded Medicaid eligibility to adults up to 138% of the Federal Poverty Level (FPL). For a single adult in 2026, 138% FPL is approximately \$1,732/month. With a monthly income of \$1,800, you appear to be just above the Expanded Medicaid income limit. [Source: Vermont Medicaid Adult Eligibility Rules, Chapter 2, Section 2.1]

MARKETPLACE SUBSIDIES:
At \$1,800/month (\$21,600 annually), you would likely qualify for ACA premium tax credits through Vermont Health Connect, as income between 100–400% FPL qualifies for subsidies...

⚠ **DISCLAIMER:** This analysis is based on current Vermont Medicaid eligibility documents but may not reflect recent changes. For an official eligibility determination, please contact Vermont Health Connect at 1-855-899-9600 or visit vermanthealthconnect.gov.

Role-Based Intelligence: Tailoring Responses to User Expertise

One of the AI Analyst's design principles is that responses should be calibrated to the user's expertise and role. A Free-tier user asking about APMs likely wants a conceptual explanation. An Advisory-tier subscriber asking about ACO REACH benchmark methodology is likely a healthcare economist who wants analytical precision.

Role-based routing:

Tier	Model	Context	Response Style
Free	Llama 3.1-8b	Standard	Accessible, definitional, educational
Subscriber	Llama 3.3-70b	Standard	Analytical, well-cited, technical depth
Student	Llama 3.3-70b	Standard	Educational, structured, scaffolded
Professional	Llama 3.3-70b	Standard	Expert-level, analytical, decision-oriented
Advisory	Claude Sonnet 4.6	Enhanced	Strategic, nuanced, advisory-quality

The model routing is not only about capability — it is about economic sustainability for the platform. Advisory subscribers who pay premium rates receive the most capable model; free users receive a capable but less powerful model.

The Advisory AI experience:

Claude Sonnet 4.6 — the AI model powering the Advisory tier — brings several capabilities that distinguish Advisory AI responses:

- **Longer reasoning chains:** More complex analytical questions receive more thorough reasoning
- **Uncertainty acknowledgment:** The model is more calibrated about acknowledging what the evidence does and does not support
- **Strategic framing:** Responses are framed in terms of organizational decisions and tradeoffs, not just information delivery
- **Vermont specificity:** Advisory users (primarily Vermont policymakers, hospital executives, and state agency staff) receive responses that integrate Vermont-specific data and context more deeply

The RAG System's Limitations: What the AI Analyst Is Not

Understanding what the AI Analyst cannot do is as important as understanding what it can:

Not a real-time news service. The knowledge base is updated when new content is published to Sanity or when HTR manually re-ingests documents. Very recent policy changes (a CMS rule published yesterday) may not be in the knowledge base yet.

Not a legal or clinical advisor. The AI Analyst provides information and analysis. It is not a licensed attorney, physician, or actuary. Decisions with significant legal, clinical, or financial consequences should be made with qualified professional advice.

Not infallible. RAG systems can retrieve incorrect or out-of-date context, and language models can misinterpret retrieved context. The AI Analyst includes citations specifically so users can verify source documents directly.

Not a substitute for human expertise in complex situations. For complex, high-stakes decisions — designing a Medicaid waiver, negotiating an ACO contract, implementing a clinical quality program — HTR's human advisory practice provides the level of expertise that no AI system can fully replicate.

Key Concepts Introduced in This Chapter

- RAG (Retrieval-Augmented Generation):** An AI architecture that combines information retrieval from a curated knowledge base with language model generation, producing answers that are grounded in retrieved evidence rather than model memorization.
- Vector embedding:** A numerical representation of text in a high-dimensional space; similar meanings are encoded as vectors that are mathematically close to each other; enables semantic search.
- pgvector:** A PostgreSQL extension enabling vector similarity search; used in HTR's Supabase database to store and query document embeddings.
- BM25:** A keyword-based document ranking algorithm (Best Match 25); used alongside vector search in hybrid retrieval to capture keyword-specific matches that semantic search may miss.
- Reciprocal Rank Fusion (RRF):** An algorithm for merging ranked lists from multiple retrieval methods; used to combine BM25 and vector search results.
- Cross-encoder reranking:** A neural reranking step that evaluates query-document relevance jointly (rather than independently); more accurate than bi-encoder similarity but computationally more expensive; used in FlashRank to rerank top candidates after hybrid retrieval.
- LLM routing:** The practice of directing user queries to different AI models based on subscription tier; balances capability and cost.