

Health Transformation Platform — Documentation

Written from scratch on **2026-09-07** by reading the actual codebase, the live Sanity dataset, and the running configuration. The code and the book are the source of truth here, not any earlier documentation.

The set

#	Document	Read it when you need to...
1	Technical Architecture	Understand what every vendor and module does and how they connect
2	Infrastructure & Maintenance	Work on the database, hosting, AI models, or CMS — or upgrade anything
3	User Guide	Understand what the product does, or perform a routine operational task

The 60-second orientation

One Next.js app on Vercel, serving four domains from a single deployment. Behind it: **Supabase** (Postgres + auth + the AI's vector store), **Sanity** (CMS, editable at [/studio](#)), and a **FastAPI service on Fly.io** that powers the AI Analyst.

Three AI vendors, doing three different jobs: **Groq** writes most chat answers, **Anthropic** handles the advisory tier, and **OpenAI** supplies embeddings, text-to-speech, and a fallback model. If OpenAI runs out of credit, chat keeps working but search quality degrades — a quiet failure worth knowing about.

A book, *Transforming Healthcare*, is the intellectual source. It lives in this repo as Markdown with its own build pipeline ([./book.sh](#)), and its framework — **five pillars held to the Equity Imperative** — drives the site's taxonomy from [frontend/lib/taxonomy/pillars.ts](#).

Related documents kept elsewhere

File	Purpose
CLAUDE.md (repo root)	Working notes and hard rules, especially for the book
BOOK_WORKFLOW.md	The author's Google Docs ↔ Markdown round-trip
HTR_ADMIN_RUNBOOK.md	Older ops runbook — useful, but verify against doc 2
book-build/README.md	The book styling/build pipeline in detail

Superseded

[frontend/docs/platform-documentation/](#) (files **00–10**) predates major changes to the platform and the book's framework. Treat it as historical. Where it conflicts with this set, this set is correct.